

Deepak Shivani

9993883717 | deepakshivani1234@gmail.com
Linkedin | Github | Open to Relocation

PROFESSIONAL SUMMARY

Data Engineer with 6 years of experience designing and operating large-scale batch and streaming data platforms on Spark, Kafka, Airflow, Databricks, and AWS. Architected lakehouse pipelines processing 500M+ records daily and led design and delivery across teams of up to 5 engineers. Strong focus on distributed-systems design, cost and performance optimization, reliability, and end-to-end production ownership.

SKILLS

Languages: Python, SQL, Java

Big Data & Processing: PySpark, Apache Spark, Kafka, Spark Streaming, Hadoop

Workflow & DevOps: Apache Airflow, Terraform, Docker, Kubernetes, Git, GitHub Actions, Bash

Cloud & Platforms: AWS (S3, RDS, Lambda, Lake Formation, Step Function, Glue), Databricks, HDFS

Other tools: FastAPI, Pandas, dbt

Data Architecture: Data Lakehouse, Medallion Architecture, Dimensional Modeling, ETL/ELT, Data Warehousing

EXPERIENCE

Adastra Group

July 2025 – Present

Senior Data Engineer

- Designed and own an end-to-end sales-funnel analytics platform integrating Google Analytics, CRM, and business datasets into KPI reporting pipelines serving top executive and core marketing teams.
- Re-architected a monolithic ETL system into modular business-layer pipeline components, cutting processing time from 7 hours to 48 minutes (monthly runs) and improving failure isolation and maintainability.
- Architected a Spark-based Bronze/Silver/Gold lakehouse processing 500M+ records daily, reducing compute costs 40% through right grain processing, Hudi optimizations and handled late arriving data with idempotency.
- Own orchestration, monitoring, infrastructure provisioning (Terraform), and CI/CD for production workloads.

Guidewire Software

June 2024 – July 2025

Data Engineer II

- Designed 15+ customer-facing datamarts supporting 500+ business users and executive reporting.
- Built Spark medallion-architecture pipelines for multi-source ingestion (6 sources has around 10 TB daily) with different SLAs (On-demand / 5 hours SLA).
- Developed orchestration workflows using AWS Step Functions, terraform and Lambda for scalable ETL execution.
- Drove an Apache Iceberg + Starburst lakehouse POC, reducing compute cost 30% via metadata filtering, predicate pushdown, separation of compute and storage.

STL Digital

October 2022 – May 2024

Data Engineer

- Led a 5-engineer team across a 6-month project, owning design and delivery of OpenAPI-spec-based configurations
- Built batch and streaming pipelines (Kafka, Spark, AWS) powering large-scale test-execution platform analytics across 25+ internal teams, handling 3 million executions daily.
- Developed PySpark reporting pipelines delivering retention, usage-pattern, and daily/weekly/monthly execution analytics for operational insight.
- Built FastAPI REST services and Airflow workflows orchestrating distributed test execution across 100 mobile devices. Contributed to an image-validation system using OpenCV.

Tata Consultancy Services

September 2020 – September 2022

Data Engineer

- Maintained ETL pipelines on Azure Data Factory and Databricks supporting CRM ingestion and reporting.
- Developed SparkSQL transformation logic and monitored scheduled batch pipelines for operational reliability.

CERTIFICATIONS & AWARDS

- Microsoft Certified: Azure Data Engineer Associate
- AWS Certified Data Engineer - Associate

EDUCATION

Rajiv Gandhi Technical University, Bhopal

Bachelor of Engineering

Aug. 2016 – Aug. 2020

Mechanical Engineering